

MICAH GOLDBLUM

micah.g@columbia.edu, goldblum.github.io

EMPLOYMENT

Columbia University, Department of Electrical Engineering Assistant Professor	<i>July 2024 - Present</i>
New York University Postdoctoral Researcher (Advised by Professors Yann LeCun and Andrew Gordon Wilson)	<i>September 2021 - June 2024</i>

EDUCATION

University of Maryland Ph.D. in Mathematics	<i>September 2014 - May 2020</i>
University of Maryland B.Sc. in Mathematics	<i>September 2010 - May 2014</i>

SELECTED HONORS AND AWARDS

Google Research Award	<i>2026</i>
Blavatnik Awards for Young Scientists Finalist	<i>2023</i>
NYU Postdoctoral Research and Professional Development Support Grant	<i>2023</i>
ICML Outstanding Paper Award	<i>2022</i>
Seymour Goldberg Gold Medal - Spotlight on Graduate Student Research	<i>2016</i>
University of Maryland Dean's Fellowship	<i>2014, 2015</i>
Phi Beta Kappa	<i>2014</i>

SELECTED PREPRINTS

A Cookbook of Self-Supervised Learning Randall Balestrierio, Mark Ibrahim, ... , Hamed Pirsiavash, Yann LeCun, Micah Goldblum
Commercial LLM Agents Are Already Vulnerable to Simple Yet Dangerous Attacks Ang Li, Yin Zhou, Vethavikashini Chithrha Raghuram, Tom Goldstein, Micah Goldblum

PUBLICATIONS

Answering At Any Cost: Frontier LLMs Are Consequence-Insensitive Arka Pal, Au Kwok Chun, Louai Zahran, Rahul Krishna Thomas, Kevin David Hayes, Adam Block, Micah Goldblum <i>Advances in Neural Information Processing (NeurIPS) 2026</i>	<i>2026</i>
Dynamic Delayed Tree Expansion For Improved Multi-Path Speculative Decoding Rahul Krishna Thomas, Teo Kitanovski, Micah Goldblum , Arka Pal <i>Advances in Neural Information Processing (NeurIPS) 2026</i>	<i>2026</i>
Understanding Reasoning from Pretraining to Post-Training: Chess as a Controlled Testbed Jingyan Shen, Ang Li, Salman Rahman, Yifan Sun, Micah Goldblum , Matus Telgarsky, Pavel Izmailov <i>Advances in Neural Information Processing (NeurIPS) 2026</i>	<i>2026</i>
Zebra-CoT: A Dataset for Interleaved Vision-Language Reasoning Ang Li, Charles L. Wang, Kaiyu Yue, Zikui Cai, Ollie Liu, Deqing Fu, ..., Micah Goldblum <i>International Conference on Learning Representations (ICLR) 2026</i>	<i>2026</i>
Identifying and Evaluating Inactive Heads in Pretrained LLMs Pedro Sandoval-Segura, Xijun Wang, Ashwinee Panda, Micah Goldblum , Ronen Basri, Tom Goldstein, David Jacobs <i>International Conference on Learning Representations (ICLR) 2026</i>	<i>2026</i>
Teaching Pretrained Language Models to Think Deeper with Retrofitted Recurrence Sean McLeish, Ang Li, John Kirchenbauer, Dayal Singh Kalra, Brian R Bartoldson, Bhavya Kailkhura, Avi Schwarzschild, Jonas Geiping, Tom Goldstein, Micah Goldblum <i>Conference on Language Modeling (COLM) 2026</i>	<i>2026</i>
LiveMathematicianBench: A Live Benchmark for Mathematician-Level Reasoning with Proof Sketches	<i>2026</i>

- Linyang He, Qiyao Yu, Hanze Dong, Baohao Liao, Xinxing Xu, **Micah Goldblum**, Jiang Bian, Nima Mesgarani
Conference on Language Modeling (COLM) 2026
- Pinocchio: Estimating the Uncertainty of Black-Box Language Models** 2026
Kevin David Hayes, Arka Pal, Haosong Zhang, Tom Goldstein, **Micah Goldblum**
Conference on Language Modeling (COLM) 2026
- Recovering Wasted Compute in Autoresearch Agents** 2026
Abhigyan Acherjee, Au Kwok Chun, Amrutha Rao, Zaiqian Chen, Kazem Meidani, Bayan Bruss, **Micah Goldblum**
Conference on Language Modeling (COLM) 2026
- Vista4D: Video Reshooting with 4D Point Clouds** 2026
Kuan Heng Lin, Zhizheng Liu, ..., Bolei Zhou, **Micah Goldblum**, Paul Debevec, Ning Yu
Computer Vision and Pattern Recognition Conference (CVPR) 2026
- Small Batch Size Training for Language Models: When Vanilla SGD Works, and Why Gradient Accumulation Is Wasteful** 2025
Martin Marek, Sanae Lotfi, Aditya Somasundaram, Andrew Gordon Wilson, **Micah Goldblum**
Advances in Neural Information Processing Systems (NeurIPS) 2025
- Gemstones: A Model Suite for Scaling Laws** 2025
Sean Michael McLeish, John Kirchenbauer, David Yu Miller, Siddharth Singh, Abhinav Bhatele, **Micah Goldblum**, Ashwinee Panda, Tom Goldstein
Advances in Neural Information Processing Systems (NeurIPS) 2025
- Brain-Predictive Reasoning Embedding through Residual Disentanglement** 2025
Linyang He, Tianjun Zhong, Richard Antonello, Gavin Mischler, **Micah Goldblum**, Nima Mesgarani
Advances in Neural Information Processing Systems (NeurIPS) 2025
- FineGRAIN: Evaluating Failure Modes of Text-to-Image Models with Vision Language Model Judges** 2025
Kevin David Hayes, **Micah Goldblum**, Vikash Sehwal, Gowthami Somepalli, Ashwinee Panda, Tom Goldstein
Advances in Neural Information Processing Systems (NeurIPS) 2025
- LiveBench: A Challenging, Contamination-Free LLM Benchmark** 2025
Colin White, Samuel Dooley, Manley Roberts, Arka Pal, Benjamin Feuer, ... , Yann LeCun, Tom Goldstein, Willie Neiswanger, **Micah Goldblum**
International Conference on Learning Representations (ICLR) 2025
- Style Outweighs Substance: Failure Modes of LLM Judges in Alignment Benchmarking** 2025
Benjamin Feuer, **Micah Goldblum**, Teresa Datta, Sanjana Nambiar, Raz Besaleli, Samuel Dooley, Max Cembalest, John P Dickerson
International Conference on Learning Representations (ICLR) 2025
- Adaptive Retention & Correction for Continual Learning** 2025
Haoran Chen, **Micah Goldblum**, Zuxuan Wu, Yu-Gang Jiang
International Conference on Learning Representations (ICLR) 2025
- Hidden No More: Attacking and Defending Private Third-Party LLM Inference** 2025
Arka Pal, Rahul Krishna Thomas, Louai Zahran, Erica Choi, Akilesh Potti, **Micah Goldblum**
International Conference on Machine Learning (ICML) 2025
- Kolmogorov Complexity, the No Free Lunch Theorem, and the Role of Inductive Biases in Machine Learning** 2024
Micah Goldblum, Marc Anton Finzi, Keefer Rowan, Andrew Gordon Wilson
International Conference on Machine Learning (ICML) 2024
- Spotting LLMs With Binoculars: Zero-Shot Detection of Machine-Generated Text** 2024
Abhimanyu Hans, Avi Schwarzschild, Valeriia Cherepanova, Hamid Kazemi, Aniruddha Saha, **Micah Goldblum**, Jonas Geiping, Tom Goldstein
International Conference on Machine Learning (ICML) 2024
- Non-vacuous generalization bounds for large language models** 2024
Sanae Lotfi, Marc Finzi, Yilun Kuang, Tim GJ Rudner,

- Micah Goldblum**, Andrew Gordon Wilson
International Conference on Machine Learning (ICML) 2024
- Large Language Models Must Be Taught to Know What They Dont Know** 2024
Sanyam Kapoor, Nate Gruver, Manley Roberts, Arka Pal, Samuel Dooley, Katherine Collins, Umang Bhatt, Adrian Weller, **Micah Goldblum**, Andrew Gordon Wilson
Advances in Neural Information Processing Systems (NeurIPS) 2024
- TuneTables: Context Optimization for Scalable Prior-Data Fitted Networks** 2024
Benjamin Feuer, Robin Tibor Schirrmeister, Valeriia Cherepanova, Chinmay Hegde, Frank Hutter, **Micah Goldblum**, Niv Cohen, Colin White
Advances in Neural Information Processing Systems (NeurIPS) 2024
- Unlocking Tokens as Data Points for Generalization Bounds on Larger Language Models** 2024
Sanae Lotfi, Yilun Kuang, Marc Anton Finzi, Brandon Amos, **Micah Goldblum**, Andrew Gordon Wilson
Advances in Neural Information Processing Systems (NeurIPS) 2024
- Searching for Efficient Linear Layers over a Continuous Space of Structured Matrices** 2024
Andres Potapczynski, Shikai Qiu, Marc Anton Finzi, Christopher Ferri, Zixi Chen, **Micah Goldblum**, C. Bayan Bruss, Christopher De Sa, Andrew Gordon Wilson
Advances in Neural Information Processing Systems (NeurIPS) 2024
- Compute Better Spent: Replacing Dense Layers with Structured Matrices** 2024
Shikai Qiu, Andres Potapczynski, Marc Anton Finzi, **Micah Goldblum**, Andrew Gordon Wilson
International Conference on Machine Learning (ICML) 2024
- Measuring Style Similarity in Diffusion Models** 2024
Gowthami Somepalli, Anubhav Gupta, Kamal Gupta, Shramay Palta, **Micah Goldblum**, Jonas Geiping, Abhinav Shrivastava, Tom Goldstein
European Conference on Computer Vision (ECCV) 2024
- On the Reliability of Watermarks for Large Language Models** 2024
John Kirchenbauer, Jonas Geiping, Yuxin Wen, Manli Shu, Khalid Saifullah, Kezhi Kong, Kasun Fernando, Aniruddha Saha, **Micah Goldblum**, Tom Goldstein
International Conference on Learning Representations (ICLR) 2024
- Universal guidance for diffusion models** 2024
Arpit Bansal, Hong-Min Chu, Avi Schwarzschild, Soumyadip Sengupta, **Micah Goldblum**, Jonas Geiping, Tom Goldstein
International Conference on Learning Representations (ICLR) 2024
- NEFTune: Noisy Embeddings Improve Instruction Finetuning** 2024
Neel Jain, Ping-yeh Chiang, Yuxin Wen, John Kirchenbauer, Hong-Min Chu, Gowthami Somepalli, Brian R Bartoldson, Bhavya Kaillkhura, Avi Schwarzschild, Aniruddha Saha, **Micah Goldblum**, Jonas Geiping, Tom Goldstein
International Conference on Learning Representations (ICLR) 2024
- Battle of the Backbones: A Large-Scale Comparison of Pretrained Models across Computer Vision Tasks** 2023
Micah Goldblum, Hossein Sourì, Renkun Ni, Manli Shu, Viraj Uday Prabhu, Gowthami Somepalli, Prithvijit Chattopadhyay, Adrien Bardes, Mark Ibrahim, Judy Hoffman, Rama Chellappa, Andrew Gordon Wilson, Tom Goldstein
Advances in Neural Information Processing Systems (NeurIPS) 2023
- Rethinking Bias Mitigation: Fairer Architectures Make for Fairer Face Recognition** 2023
Samuel Dooley, Rhea Sukthanker, John P Dickerson, Colin White, Frank Hutter, **Micah Goldblum**
Advances in Neural Information Processing Systems (NeurIPS) 2023
- Transfer Learning with Deep Tabular Models** 2023
Roman Levin, Valeriia Cherepanova, Avi Schwarzschild, Arpit Bansal,

C. Bayan Bruss, Tom Goldstein, Andrew Gordon Wilson, **Micah Goldblum**
International Conference on Learning Representations (ICLR) 2023

Simplifying Neural Network Training Under Class Imbalance 2023
 Ravid Shwartz-Ziv*, **Micah Goldblum***, Yucen Lily Li, C. Bayan Bruss, Andrew Gordon Wilson
Advances in Neural Information Processing Systems (NeurIPS) 2023

A Performance-Driven Benchmark for Feature Selection in Tabular Deep Learning 2023
 Valeriia Cherepanova, Gowthami Somepalli, Jonas Geiping, C. Bayan Bruss,
 Andrew Gordon Wilson, Tom Goldstein, **Micah Goldblum**
Advances in Neural Information Processing Systems (NeurIPS) 2023

Gradient-based optimization is not necessary for generalization in neural networks 2023
 Ping-yeh Chiang, Renkun Ni, David Yu Miller, Arpit Bansal, Jonas Geiping,
Micah Goldblum, Tom Goldstein
International Conference on Learning Representations (ICLR) 2023

Cold Diffusion: Inverting Arbitrary Image Transforms Without Noise 2023
 Arpit Bansal, Eitan Borgnia, Hong-Min Chu, Jie S. Li, Hamid Kazemi, Furong Huang,
Micah Goldblum, Jonas Geiping, Tom Goldstein
Advances in Neural Information Processing Systems (NeurIPS) 2023

Why Diffusion Models Memorize and How to Mitigate Copying 2023
 Gowthami Somepalli, Vasu Singla, **Micah Goldblum**,
 Jonas Geiping, Tom Goldstein
Advances in Neural Information Processing Systems (NeurIPS) 2023

Hard Prompts Made Easy: Gradient-Based Discrete Optimization for Prompt Tuning and Discovery 2023
 Yuxin Wen, Neel Jain, John Kirchenbauer, **Micah Goldblum**, Jonas Geiping, Tom Goldstein
Advances in Neural Information Processing Systems (NeurIPS) 2023

When Do Neural Nets Outperform Boosted Trees on Tabular Data? 2023
 Duncan C. McElfresh, Sujay Khandagale, Jonathan Valverde, Vishak Prasad C,
 Ganesh Ramakrishnan, **Micah Goldblum**, Colin White
Advances in Neural Information Processing Systems (NeurIPS) 2023

What Can We Learn from Unlearnable Datasets? 2023
 Pedro Sandoval-Segura, Vasu Singla, Jonas Geiping, **Micah Goldblum**, Tom Goldstein
Advances in Neural Information Processing Systems (NeurIPS) 2023

Diffusion Art or Digital Forgery? Investigating Data Replication in Diffusion Models 2023
 Gowthami Somepalli, Vasu Singla, **Micah Goldblum**, Jonas Geiping, Tom Goldstein
Computer Vision and Pattern Recognition Conference (CVPR) 2023

The Lie Derivative for Measuring Learned Equivariance 2023
 Nate Gruver, Marc Anton Finzi, **Micah Goldblum**, Andrew Gordon Wilson
International Conference on Learning Representations (ICLR) 2023

Exploring and Exploiting Decision Boundary Dynamics for Adversarial Robustness 2023
 Yuancheng Xu, Yanchao Sun, **Micah Goldblum**, Tom Goldstein, Furong Huang
International Conference on Learning Representations (ICLR) 2023

Canary in a Coalmine: Better Membership Inference with Ensembled Adversarial Queries 2023
 Yuxin Wen, Arpit Bansal, Hamid Kazemi, Eitan Borgnia,
Micah Goldblum, Jonas Geiping, Tom Goldstein
International Conference on Learning Representations (ICLR) 2023

Panning for Gold in Federated Learning: Targeted Text Extraction under Arbitrarily Large-Scale Aggregation 2023
 Hong-Min Chu, Jonas Geiping, Liam Fowl, **Micah Goldblum**, Tom Goldstein
International Conference on Learning Representations (ICLR) 2023

How Much Data Are Augmentations Worth? An Investigation into Scaling Laws, Invariance, and Implicit Regularization 2023
 Jonas Geiping, **Micah Goldblum**, Gowthami Somepalli, Ravid Shwartz-Ziv,

Tom Goldstein, Andrew Gordon Wilson

International Conference on Learning Representations (ICLR) 2023

Decepticons: Corrupted Transformers Breach Privacy in Federated Learning for Language Models

2023

Liam Fowl, Jonas Geiping, Steven Reich, Yuxin Wen, Wojciech Czaja, Micah Goldblum, Tom Goldstein

International Conference on Learning Representations (ICLR) 2023

A Deep Dive into Dataset Imbalance and Bias in Face Identification

2023

Valeriia Cherepanova, Steven Reich, Samuel Dooley, Hossein Sourì, Micah Goldblum, Tom Goldstein

Proceedings of the 2023 AAAI/ACM Conference on AI, Ethics, and Society (AIES)

Bayesian Model Selection, the Marginal Likelihood, and Generalization

2022

Sanae Lotfi, Pavel Izmailov, Gregory Benton, Micah Goldblum, Andrew Gordon Wilson

International Conference on Machine Learning (ICML) 2022 (Outstanding Paper Award)

Pre-Train Your Loss: Easy Bayesian Transfer Learning with Informative Priors

2022

Ravid Shwartz-Ziv*, Micah Goldblum*, Hossein Sourì, Sanyam Kapoor, Chen Zhu, Yann LeCun, Andrew Gordon Wilson

Advances in Neural Information Processing Systems (NeurIPS) 2022

Scalable Algorithm Synthesis with Recurrent Networks: Extrapolation without Overthinking

2022

Arpit Bansal, Avi Schwarzschild, Eitan Borgnia, Zeyad Emam, Furong Huang, Micah Goldblum, Tom Goldstein

Advances in Neural Information Processing Systems (NeurIPS) 2022

PAC-Bayes Compression Bounds So Tight That They Can Explain Generalization

2022

Sanae Lotfi, Sanyam Kapoor, Marc Anton Finzi, Andres Potapczynski, Micah Goldblum, Andrew Gordon Wilson

Advances in Neural Information Processing Systems (NeurIPS) 2022

Where do Models go Wrong? Parameter-Space Saliency Maps for Explainability

2022

Roman Levin, Manli Shu, Eitan Borgnia, Furong Huang, Micah Goldblum, Tom Goldstein

Advances in Neural Information Processing Systems (NeurIPS) 2022

Sleeper Agent: Scalable Hidden Trigger Backdoors for Neural Networks Trained from Scratch

2022

Hossein Sourì, Liam H Fowl, Rama Chellappa, Micah Goldblum, Tom Goldstein

Advances in Neural Information Processing Systems (NeurIPS) 2022

Chroma-VAE: Mitigating Shortcut Learning with Generative Classifiers

2022

Wanqian Yang, Polina Kirichenko, Micah Goldblum, Andrew Gordon Wilson

Advances in Neural Information Processing Systems (NeurIPS) 2022

Autoregressive Perturbations for Data Poisoning

2022

Pedro Sandoval-Segura, Vasu Singla, Jonas Geiping, Micah Goldblum, Tom Goldstein, David W. Jacobs

Advances in Neural Information Processing Systems (NeurIPS) 2022

Fishing for User Data in Large-Batch Federated Learning via Gradient Magnification

2022

Yuxin Wen, Jonas Geiping, Liam Fowl, Micah Goldblum, Tom Goldstein

International Conference on Machine Learning (ICML) 2022

Plug-In Inversion: Model-Agnostic Inversion for Vision with Data Augmentations

2022

Amin Ghiasi, Hamid Kazemi, Steven Reich, Chen Zhu, Micah Goldblum, Tom Goldstein

International Conference on Machine Learning (ICML) 2022

Dataset Security for Machine Learning: Data Poisoning, Backdoor Attacks, and Defenses

2022

Micah Goldblum, Dimitris Tsipras, Chulin Xie, Xinyun Chen, Avi Schwarzschild, Dawn Song, Aleksander Madry, Bo Li, Tom Goldstein

IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI) 2022

Can You Learn the Same Model Twice? Investigating Reproducibility and Double Descent from the Decision Boundary Perspective

2022

Gowthami Somepalli, Liam Fowl, Arpit Bansal, Ping Yeh-Chiang, Yehuda Dar, Richard Baraniuk,
Micah Goldblum, Tom Goldstein

Conference on Computer Vision and Pattern Recognition (CVPR) 2022

Contrastive Learning is Just Meta-Learning

2022

Renkun Ni, Manli Shu, Hossein Souri, **Micah Goldblum**, Tom Goldstein

International Conference on Learning Representations (ICLR) 2022

Stochastic Training is Not Necessary for Generalization

2022

Jonas Geiping, **Micah Goldblum**, Phil Pope, Michael Moeller, Tom Goldstein

International Conference on Learning Representations (ICLR) 2022

Robbing the Fed: Directly Obtaining Private Data in Federated Learning with Modified Models

2022

Liam H Fowl, Jonas Geiping, Wojciech Czaja, **Micah Goldblum**, Tom Goldstein

International Conference on Learning Representations (ICLR) 2022

The Uncanny Similarity of Recurrence and Depth

2022

Avi Schwarzschild, Arjun Gupta, Amin Ghiasi, **Micah Goldblum**, Tom Goldstein

International Conference on Learning Representations (ICLR) 2022

Towards Transferable Adversarial Attacks on Vision Transformers

2022

Zhipeng Wei, Jingjing Chen, **Micah Goldblum**, Zuxuan Wu, Tom Goldstein, Yu-Gang Jiang

AAAI Conference on Artificial Intelligence (AAAI) 2022

Adversarial Examples Make Strong Poisons

2021

Liam Fowl*, **Micah Goldblum***, Ping-yeh Chiang, Jonas Geiping, Wojtek Czaja, Tom Goldstein

Advances in Neural Information Processing Systems (NeurIPS) 2021

Can You Learn an Algorithm? Generalizing from Easy to Hard Problems with Recurrent Networks

2021

Avi Schwarzschild, Eitan Borgnia, Arjun Gupta, Furong Huang, Uzi Vishkin,

Micah Goldblum, Tom Goldstein

Advances in Neural Information Processing Systems (NeurIPS) 2021

Prepare for the Worst: Generalizing across Domain Shifts with Adversarial Batch Normalization

2021

Manli Shu, Zuxuan Wu, **Micah Goldblum**, Tom Goldstein

Advances in Neural Information Processing Systems (NeurIPS) 2021

Adversarial Attacks on Machine Learning Systems for High-Frequency Trading

2021

Micah Goldblum, Avi Schwarzschild, Naftali Cohen, Tucker Balch, Ankit B. Patel,
Tom Goldstein

ACM International Conference on AI in Finance (ICAIF) 2021

Just How Toxic is Data Poisoning? A Unified Benchmark for Backdoor and Data Poisoning Attacks

2021

Avi Schwarzschild*, **Micah Goldblum***, Arjun Gupta, John P Dickerson, Tom Goldstein

International Conference on Machine Learning (ICML) 2021

Data Augmentation for Meta-Learning

2021

Renkun Ni, **Micah Goldblum**, Amr Sharaf, Kezhi Kong, Tom Goldstein

International Conference on Machine Learning (ICML) 2021

The Intrinsic Dimension of Images and Its Impact on Learning

2021

Phil Pope, Chen Zhu, Ahmed Abdelkader, **Micah Goldblum**, Tom Goldstein

International Conference on Learning Representations (ICLR) 2021

LowKey: Leveraging Adversarial Attacks to Protect Social Media Users from Facial Recognition

2021

Valeriia Cherepanova, **Micah Goldblum**, Harrison Foley, Shiyuan Duan,

John P Dickerson, Gavin Taylor, Tom Goldstein

International Conference on Learning Representations (ICLR) 2021

Strong Data Augmentation Sanitizes Poisoning and Backdoor Attacks Without an Accuracy Tradeoff

2021

Eitan Borgnia, Valeriia Cherepanova, Liam Fowl, Amin Ghiasi, Jonas Geiping,

Micah Goldblum, Tom Goldstein, Arjun Gupta
International Conference on Acoustics, Speech, and Signal Processing (ICASSP) 2021

Robust Few-Shot Learning: A Meta-Learning Approach 2020
Micah Goldblum, Liam Fowl, Tom Goldstein
Advances in Neural Information Processing Systems (NeurIPS) 2020

Unraveling Meta-Learning: Understanding Feature Representations for Few-Shot Tasks 2020
Micah Goldblum, Steven Reich, Liam Fowl, Renkun Ni, Valeriia Cherepanova, Tom Goldstein
International Conference on Machine Learning (ICML) 2020

Truth or backpropaganda? An empirical investigation of deep learning theory 2020
Micah Goldblum, Jonas Geiping, Avi Schwarzschild, Michael Moeller, Tom Goldstein
International Conference on Learning Representations (ICLR) 2020

WITCHcraft: Efficient PGD attacks with random step size 2020
Ping-Yeh Chiang, Jonas Geiping, Micah Goldblum, Tom Goldstein, Renkun Ni, Steven Reich, Ali Shafahi
International Conference on Acoustics, Speech, and Signal Processing (ICASSP) 2020

Adversarially Robust Distillation 2020
Micah Goldblum, Liam Fowl, Soheil Feizi, Tom Goldstein
Proceedings of the AAAI Conference on Artificial Intelligence. Vol. 34.

SELECT INVITED TALKS

Ecole Polytechnique, Spring Masterclass (previous invited speaker was Michael I Jordan) 2026
Jump Trading, AI Symposium Keynote Speaker 2025
Columbia University, Distinguished AI Lectures 2025
Princeton University, Algorithms and Machine Learning Seminar 2025
Columbia University, Statistics Seminar 2025
Two Sigma 2025
Columbia University, Conference on AI Agents for Work 2025
Cornell University, Department of Computer Science 2024
University of Pennsylvania, Department of Computer Science 2024
Georgia Institute of Technology, Department of Interactive Computing 2024
Northeastern University, Department of Computer Science 2024
University of Edinburgh, School of Informatics 2024
University of North Carolina Chapel Hill, Department of Computer Science 2024
University of British Columbia, Department of Computer Science 2024
Arthur Panel on AI and Science 2024
Chalmers AI Research Center Workshop for Structured Learning 2023
UCLA + MPI MIS Math Machine Learning Seminar 2023
Vanderbilt Machine Learning Seminar 2023
CISPA Helmholtz Center for Information Security 2023
University of California Los Angeles, Department of Mathematics 2023
ML Collective 2023
École Polytechnique Fédérale de Lausanne (EPFL) 2022
C3.ai Digital Transformation Institute: Open Problems in Robustness 2020

MEDIA COVERAGE

Context compression finally works in production 2026
VentureBeat

How close is AI to human-level intelligence? 2024
Nature

LiveBench is an open LLM benchmark that uses contamination-free test data and objective scoring 2024
VentureBeat

A New AI Detection Tool May Have Solved False Positives for Student Writing 2024
Business Insider

AI Spam Threatens the Internet – AI Can Also Protect It 2024
IEEE Spectrum

Researchers develop new LiveBench benchmark for measuring AI models response accuracy 2024
Silicon Angle

Artists Are Slipping Anti-AI Poison into Their Art. Here's How It Works 2024
Scientific American

Binoculars is the most powerful AI text detector with over 90% accuracy 2024
The Decoder

New AI detection tool measures how “surprising” word choices are 2024
Freethink

AI and the future of work 2023
Daily Maverick

Image-generating AI can copy and paste from training data, raising IP concerns 2022
TechCrunch

How I Lost Control Over My Own Face 2021
Der Spiegel Magazine

Cómo evitar que los sistemas de reconocimiento facial descifren las fotos de tus redes 2021
El País

LowKey cool: This web app will tweak your photos to flummox facial-recognition systems, apparently 2021
The Register

COMMUNITY SERVICE

- Chair of the organizing committee for the NeurIPS 2020 Workshop on Dataset Curation and Security.
- Organizer of the NeurIPS 2023 Workshop on Backdoors in Deep Learning,
- Organizer of the NeurIPS 2024 Workshop on Red Teaming GenAI: What Can We Learn from Adversaries?
- Organizer of the NeurIPS 2024 Workshop on Scientific Methods for Understanding Neural Networks: Discovering, Validating, and Falsifying Theories of Deep Learning with Experiments.
- Chair of the organizing committee for the ICLR 2025 workshop on Building Trust in LLMs and LLM Applications: From Guardrails to Explainability to Regulation
- Served as an Area Chair, Session Chair, or Reviewer for conferences and journals including NeurIPS, ICML, ICLR, CVPR, ECCV, COLM, and TPAMI.
- Member of the steering committee for Columbia University Foundations of Data Science Center